

# C2W2\_Assignment

July 14, 2021

## 1 Breast Cancer Prediction

In this exercise, you will train a neural network on the [Breast Cancer Dataset](#) to predict if the tumor is malignant or benign.

If you get stuck, we recommend that you review the ungraded labs for this week.

### 1.1 Imports

```
[1]: import tensorflow as tf
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, Input

import numpy as np
import matplotlib.pyplot as plt
import matplotlib.ticker as mticker
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix
import itertools
from tqdm import tqdm
import tensorflow_datasets as tfds

tf.get_logger().setLevel('ERROR')
```

### 1.2 Load and Preprocess the Dataset

We first load the dataset and create a data frame using pandas. We explicitly specify the column names because the CSV file does not have column headers.

```
[2]: data_file = './data/data.csv'
col_names = ["id", "clump_thickness", "un_cell_size", "un_cell_shape",
↪ "marginal_adheshion", "single_eph_cell_size", "bare_nuclei",
↪ "bland_chromatin", "normal_nucleoli", "mitoses", "class"]
df = pd.read_csv(data_file, names=col_names, header=None)
```

```
[3]: df.head()
```

```
[3]:      id  clump_thickness  un_cell_size  un_cell_shape  marginal_adhesion  \
0  1000025                5              1              1                  1
1  1002945                5              4              4                  5
2  1015425                3              1              1                  1
3  1016277                6              8              8                  1
4  1017023                4              1              1                  3

      single_eph_cell_size  bare_nuclei  bland_chromatin  normal_nucleoli  \
0                      2              1              3              1
1                      7             10              3              2
2                      2              2              3              1
3                      3              4              3              7
4                      2              1              3              1

      mitoses  class
0          1      2
1          1      2
2          1      2
3          1      2
4          1      2
```

We have to do some preprocessing on the data. We first pop the id column since it is of no use for our problem at hand.

```
[4]: df.pop("id")
```

```
[4]: 0      1000025
1      1002945
2      1015425
3      1016277
4      1017023
...
694    776715
695    841769
696    888820
697    897471
698    897471
Name: id, Length: 699, dtype: int64
```


Upon inspection of data, you can see that some values of the **bare\_nuclei** column are unknown. We drop the rows with these unknown values. We also convert the **bare\_nuclei** column to numeric. This is required for training the model.

```
[5]: df = df[df["bare_nuclei"] != '?']
df.bare_nuclei = pd.to_numeric(df.bare_nuclei)
```

We check the class distribution of the data. You can see that there are two classes, 2.0 and 4.0  
According to the dataset: \* **2.0 = benign** \* **4.0 = malignant**

```
[6]: df['class'].hist(bins=20)
```

```
[6]: <matplotlib.axes._subplots.AxesSubplot at 0x7feeaea1c190>
```



We are going to model this problem as a binary classification problem which detects whether the tumor is malignant or not. Hence, we change the dataset so that: \* **benign(2.0) = 0** \* **malignant(4.0) = 1**

```
[7]: df['class'] = np.where(df['class'] == 2, 0, 1)
```

We then split the dataset into training and testing sets. Since the number of samples is small, we will perform validation on the test set.

```
[8]: train, test = train_test_split(df, test_size = 0.2)
```

We get the statistics for training. We can look at statistics to get an idea about the distribution of plots. If you need more visualization, you can create additional data plots. We will also be using the mean and standard deviation from statistics for normalizing the data

```
[11]: train_stats = train.describe()  
train_stats.pop('class')  
train_stats = train_stats.transpose()
```

```
[12]: train_stats
```

```
[12]:
```

	count	mean	std	min	25%	50%	75%	max
clump_thickness	546.0	4.408425	2.814745	1.0	2.0	4.0	6.00	10.0
un_cell_size	546.0	3.108059	3.039986	1.0	1.0	1.0	4.00	10.0
un_cell_shape	546.0	3.183150	3.003877	1.0	1.0	1.0	5.00	10.0
marginal_adheshion	546.0	2.800366	2.875789	1.0	1.0	1.0	3.75	10.0
single_eph_cell_size	546.0	3.214286	2.209209	1.0	2.0	2.0	4.00	10.0
bare_nuclei	546.0	3.543956	3.655195	1.0	1.0	1.0	6.00	10.0
bland_chromatin	546.0	3.421245	2.401496	1.0	2.0	3.0	4.00	10.0
normal_nucleoli	546.0	2.882784	3.086074	1.0	1.0	1.0	4.00	10.0
mitoses	546.0	1.575092	1.717381	1.0	1.0	1.0	1.00	10.0

We pop the class column from the training and test sets to create train and test outputs.

```
[13]: train_Y = train.pop("class")
      test_Y = test.pop("class")
```

Here we normalize the data by using the formula:  $X = (X - \text{mean}(X)) / \text{StandardDeviation}(X)$

```
[14]: def norm(x):
      return (x - train_stats['mean']) / train_stats['std']
```

```
[17]: norm_train_X = norm(train)
      norm_test_X = norm(test)
```

We now create Tensorflow datasets for training and test sets to easily be able to build and manage an input pipeline for our model.

```
[19]: train_dataset = tf.data.Dataset.from_tensor_slices((norm_train_X.values,
    ↪train_Y.values))
      test_dataset = tf.data.Dataset.from_tensor_slices((norm_test_X.values, test_Y.
    ↪values))
```

We shuffle and prepare a batched dataset to be used for training in our custom training loop.

```
[21]: batch_size = 32
      train_dataset = train_dataset.shuffle(buffer_size=len(train)).batch(batch_size)

      test_dataset = test_dataset.batch(batch_size=batch_size)
```

```
[22]: a = enumerate(train_dataset)

      print(len(list(a)))
```

### 1.3 Define the Model

Now we will define the model. Here, we use the Keras Functional API to create a simple network of two `Dense` layers. We have modelled the problem as a binary classification problem and hence we add a single layer with sigmoid activation as the final layer of the model.

```
[23]: def base_model():
        inputs = tf.keras.layers.Input(shape=(len(train.columns)))

        x = tf.keras.layers.Dense(128, activation='relu')(inputs)
        x = tf.keras.layers.Dense(64, activation='relu')(x)
        outputs = tf.keras.layers.Dense(1, activation='sigmoid')(x)
        model = tf.keras.Model(inputs=inputs, outputs=outputs)
        return model

model = base_model()
```

### 1.4 Define Optimizer and Loss

We use RMSprop optimizer and binary crossentropy as our loss function.

```
[24]: optimizer = tf.keras.optimizers.RMSprop(learning_rate=0.001)
        loss_object = tf.keras.losses.BinaryCrossentropy()
```

### 1.5 Evaluate Untrained Model

We calculate the loss on the model before training begins.

```
[25]: outputs = model(norm_test_X.values)
        loss_value = loss_object(y_true=test_Y.values, y_pred=outputs)
        print("Loss before training %.4f" % loss_value.numpy())
```

Loss before training 0.7284

We also plot the confusion matrix to visualize the true outputs against the outputs predicted by the model.

```
[27]: def plot_confusion_matrix(y_true, y_pred, title='', labels=[0,1]):
        cm = confusion_matrix(y_true, y_pred)
        fig = plt.figure()
        ax = fig.add_subplot(111)
        cax = ax.matshow(cm)
        plt.title(title)
        fig.colorbar(cax)
        ax.set_xticklabels([''] + labels)
        ax.set_yticklabels([''] + labels)
        plt.xlabel('Predicted')
```

```


plt.ylabel('True')
fmt = 'd'
thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, format(cm[i, j], fmt),
             horizontalalignment="center",
             color="black" if cm[i, j] > thresh else "white")
plt.show()

```

```

[28]: plot_confusion_matrix(test_Y.values, tf.round(outputs), title='Confusion Matrix_
      ↪for Untrained Model')

```



## 1.6 Define Metrics (Please complete this section)

### 1.6.1 Define Custom F1Score Metric

In this example, we will define a custom F1Score metric using the formula.

$$\text{F1 Score} = 2 * ((\text{precision} * \text{recall}) / (\text{precision} + \text{recall}))$$

$$\text{precision} = \text{true\_positives} / (\text{true\_positives} + \text{false\_positives})$$

$$\text{recall} = \text{true\_positives} / (\text{true\_positives} + \text{false\_negatives})$$

We use `confusion_matrix` defined in `tf.math` to calculate precision and recall.

Here you can see that we have subclassed `tf.keras.Metric` and implemented the three required methods `update_state`, `result` and `reset_states`.

### 1.6.2 Please complete the `result()` method:

```
[29]: class F1Score(tf.keras.metrics.Metric):

    def __init__(self, name='f1_score', **kwargs):
        '''initializes attributes of the class'''

        # call the parent class init
        super(F1Score, self).__init__(name=name, **kwargs)

        # Initialize Required variables
        # true positives
        self.tp = tf.Variable(0, dtype = 'int32')
        # false positives
        self.fp = tf.Variable(0, dtype = 'int32')
        # true negatives
        self.tn = tf.Variable(0, dtype = 'int32')
        # false negatives
        self.fn = tf.Variable(0, dtype = 'int32')

    def update_state(self, y_true, y_pred, sample_weight=None):
        '''
        Accumulates statistics for the metric

        Args:
            y_true: target values from the test data
            y_pred: predicted values by the model
        '''

        # Calulcate confusion matrix.
        conf_matrix = tf.math.confusion_matrix(y_true, y_pred, num_classes=2)

        # Update values of true positives, true negatives, false positives and
        ↪ false negatives from confusion matrix.
        self.tn.assign_add(conf_matrix[0][0])
        self.tp.assign_add(conf_matrix[1][1])
        self.fp.assign_add(conf_matrix[0][1])
        self.fn.assign_add(conf_matrix[1][0])

    def result(self):
        '''Computes and returns the metric value tensor.'''

        # Calculate precision
```

```

    if (self.tp + self.fp == 0):
        precision = 1.0
    else:
        precision = self.tp / (self.tp + self.fp)

    # Calculate recall
    if (self.tp + self.fn == 0):
        recall = 1.0
    else:
        recall = self.tp / (self.tp + self.fn)

    # Return F1 Score
    ### START CODE HERE ###
    f1_score = 2 * ((precision * recall) / (precision + recall))
    ### END CODE HERE ###

    return f1_score

def reset_states(self):
    '''Resets all of the metric state variables.'''

    # The state of the metric will be reset at the start of each epoch.
    self.tp.assign(0)
    self.tn.assign(0)
    self.fp.assign(0)
    self.fn.assign(0)

```

[30]: *# Test Code:*

```

test_F1Score = F1Score()

test_F1Score.tp = tf.Variable(2, dtype = 'int32')
test_F1Score.fp = tf.Variable(5, dtype = 'int32')
test_F1Score.tn = tf.Variable(7, dtype = 'int32')
test_F1Score.fn = tf.Variable(9, dtype = 'int32')
test_F1Score.result()

```

[30]: <tf.Tensor: shape=(), dtype=float64, numpy=0.2222222222222222>

### Expected Output:

<tf.Tensor: shape=(), dtype=float64, numpy=0.2222222222222222>

We initialize the separate metrics required for training and validation. In addition to our custom F1Score metric, we are also using BinaryAccuracy defined in `tf.keras.metrics`

[31]: `train_f1score_metric = F1Score()`  
`val_f1score_metric = F1Score()`



```
train_acc_metric = tf.keras.metrics.BinaryAccuracy()
val_acc_metric = tf.keras.metrics.BinaryAccuracy()
```

## 1.7 Apply Gradients (Please complete this section)

The core of training is using the model to calculate the logits on specific set of inputs and compute the loss(in this case **binary crossentropy**) by comparing the predicted outputs to the true outputs. We then update the trainable weights using the optimizer algorithm chosen. The optimizer algorithm requires our computed loss and partial derivatives of loss with respect to each of the trainable weights to make updates to the same.

We use gradient tape to calculate the gradients and then update the model trainable weights using the optimizer.

### 1.7.1 Please complete the following function:

```
[32]: def apply_gradient(optimizer, loss_object, model, x, y):
    '''
    applies the gradients to the trainable model weights

    Args:
        optimizer: optimizer to update model weights
        loss_object: type of loss to measure during training
        model: the model we are training
        x: input data to the model
        y: target values for each input
    '''

    with tf.GradientTape() as tape:
        ### START CODE HERE ###
        logits = model(x)
        loss_value = loss_object(y_true=y, y_pred=logits)

        gradients = tape.gradient(loss_value, model.trainable_weights)
        optimizer.apply_gradients(zip(gradients, model.trainable_weights))
        ### END CODE HERE ###

    return logits, loss_value

[33]: # Test Code:

test_model = tf.keras.models.load_model('./test_model')
test_logits, test_loss = apply_gradient(optimizer, loss_object, test_model,
    ↪ norm_test_X.values, test_Y.values)
```

```

print(test_logits.numpy()[:8])
print(test_loss.numpy())

del test_model
del test_logits
del test_loss

```

```

[[0.52413803]
 [0.5602368 ]
 [0.549544  ]
 [0.54723155]
 [0.48679572]
 [0.5337864 ]
 [0.4776442 ]
 [0.5239992 ]]
0.7068926

```

### Expected Output:

The output will be close to these values:

```

[[0.5516499 ]
 [0.52124363]
 [0.5412698 ]
 [0.54203206]
 [0.50022954]
 [0.5459626 ]
 [0.47841492]
 [0.54381996]]
0.7030578

```

## 1.8 Training Loop (Please complete this section)

This function performs training during one epoch. We run through all batches of training data in each epoch to make updates to trainable weights using our previous function. You can see that we also call `update_state` on our metrics to accumulate the value of our metrics.

We are displaying a progress bar to indicate completion of training in each epoch. Here we use `tqdm` for displaying the progress bar.

### 1.8.1 Please complete the following function:

```

[34]: def train_data_for_one_epoch(train_dataset, optimizer, loss_object, model,
                                     train_acc_metric, train_f1score_metric,
                                     verbose=True):
    """

```

*Computes the loss then updates the weights and metrics for one epoch.*

*Args:*

*train\_dataset: the training dataset*  
*optimizer: optimizer to update model weights*  
*loss\_object: type of loss to measure during training*  
*model: the model we are training*  
*train\_acc\_metric: calculates how often predictions match labels*  
*train\_f1score\_metric: custom metric we defined earlier*

*'''*

losses = []

*#Iterate through all batches of training data*

**for** step, (x\_batch\_train, y\_batch\_train) **in** enumerate(train\_dataset):

*#Calculate loss and update trainable variables using optimizer*

*### START CODE HERE ###*

logits, loss\_value = apply\_gradient(optimizer, loss\_object, model,   
↪x\_batch\_train, y\_batch\_train)

losses.append(loss\_value)

*### END CODE HERE ###*

*#Round off logits to nearest integer and cast to integer for calculating*  
↪*metrics*

logits = tf.round(logits)

logits = tf.cast(logits, 'int64')

*#Update the training metrics*

*### START CODE HERE ###*

train\_acc\_metric.update\_state(y\_true=y\_batch\_train, y\_pred=logits)

train\_f1score\_metric.update\_state(y\_true=y\_batch\_train, y\_pred=logits)

*### END CODE HERE ###*

*#Update progress*

**if** verbose:

print("Training loss for step %s: %.4f" % (int(step),   
↪float(loss\_value)))

**return** losses

[36]: *# TEST CODE*

test\_model = tf.keras.models.load\_model('./test\_model')

test\_losses = train\_data\_for\_one\_epoch(train\_dataset, optimizer, loss\_object,   
↪test\_model,

```

train_acc_metric, train_f1score_metric,
↪ verbose=False)

for test_loss in test_losses:
    print(test_loss.numpy())

del test_model
del test_losses

```

```

0.7552004
0.6243548
0.5939014
0.49887848
0.47701666
0.40541014
0.36220235
0.40997204
0.3567308
0.29308903
0.30118644
0.27259445
0.35707676
0.2417258
0.20194711
0.2440404
0.17027119
0.19259492

```

### Expected Output:

The losses should generally be decreasing and will start from around 0.75. For example:

```

0.7600615
0.6092045
0.5525634
0.4358902
0.4765755
0.43327087
0.40585428
0.32855004
0.35755336
0.3651728
0.33971977
0.27372319
0.25026917
0.29229593
0.242178
0.20602849

```

0.15887335  
0.090397514

At the end of each epoch, we have to validate the model on the test dataset. The following function calculates the loss on test dataset and updates the states of the validation metrics.

```
[37]: def perform_validation():  
    losses = []  
  
    #Iterate through all batches of validation data.  
    for x_val, y_val in test_dataset:  
  
        #Calculate validation loss for current batch.  
        val_logits = model(x_val)  
        val_loss = loss_object(y_true=y_val, y_pred=val_logits)  
        losses.append(val_loss)  
  
        #Round off and cast outputs to either 0 or 1  
        val_logits = tf.cast(tf.round(model(x_val)), 'int64')  
  
        #Update validation metrics  
        val_acc_metric.update_state(y_val, val_logits)  
        val_f1score_metric.update_state(y_val, val_logits)  
  
    return losses
```

Next we define the training loop that runs through the training samples repeatedly over a fixed number of epochs. Here we combine the functions we built earlier to establish the following flow: 1. Perform training over all batches of training data. 2. Get values of metrics. 3. Perform validation to calculate loss and update validation metrics on test data. 4. Reset the metrics at the end of epoch. 5. Display statistics at the end of each epoch.

**Note :** We also calculate the training and validation losses for the whole epoch at the end of the epoch.

```
[38]: # Iterate over epochs.  
epochs = 5  
epochs_val_losses, epochs_train_losses = [], []  
  
for epoch in range(epochs):  
    print('Start of epoch %d' % (epoch,))  
    #Perform Training over all batches of train data  
    losses_train = train_data_for_one_epoch(train_dataset, optimizer,  
→loss_object, model, train_acc_metric, train_f1score_metric)  
  
    # Get results from training metrics  
    train_acc = train_acc_metric.result()  
    train_f1score = train_f1score_metric.result()
```

```

#Perform validation on all batches of test data
losses_val = perform_validation()

# Get results from validation metrics
val_acc = val_acc_metric.result()
val_f1score = val_f1score_metric.result()

#Calculate training and validation losses for current epoch
losses_train_mean = np.mean(losses_train)
losses_val_mean = np.mean(losses_val)
epochs_val_losses.append(losses_val_mean)
epochs_train_losses.append(losses_train_mean)

print('\n Epcoh %s: Train loss: %.4f  Validation Loss: %.4f, Train Accuracy:
↪ %.4f, Validation Accuracy %.4f, Train F1 Score: %.4f, Validation F1 Score:↪
↪%.4f' % (epoch, float(losses_train_mean), float(losses_val_mean),↪
↪float(train_acc), float(val_acc), train_f1score, val_f1score))

#Reset states of all metrics
train_acc_metric.reset_states()
val_acc_metric.reset_states()
val_f1score_metric.reset_states()
train_f1score_metric.reset_states()

```

Start of epoch 0

```

Training loss for step 0: 0.7127
Training loss for step 1: 0.5764
Training loss for step 2: 0.4708
Training loss for step 3: 0.4578
Training loss for step 4: 0.4218
Training loss for step 5: 0.3714
Training loss for step 6: 0.3532
Training loss for step 7: 0.3666
Training loss for step 8: 0.2898
Training loss for step 9: 0.3062
Training loss for step 10: 0.2359
Training loss for step 11: 0.1876
Training loss for step 12: 0.2370
Training loss for step 13: 0.1951
Training loss for step 14: 0.2294
Training loss for step 15: 0.1303
Training loss for step 16: 0.1295
Training loss for step 17: 0.1334

```

```

Epcoh 0: Train loss: 0.3225  Validation Loss: 0.1413, Train Accuracy: 0.9074,
Validation Accuracy 0.9750, Train F1 Score: 0.8816, Validation F1 Score: 0.9600
Start of epoch 1

```

Training loss for step 0: 0.1034  
Training loss for step 1: 0.2242  
Training loss for step 2: 0.1379  
Training loss for step 3: 0.1278  
Training loss for step 4: 0.1505  
Training loss for step 5: 0.0914  
Training loss for step 6: 0.1202  
Training loss for step 7: 0.1210  
Training loss for step 8: 0.1361  
Training loss for step 9: 0.0579  
Training loss for step 10: 0.0493  
Training loss for step 11: 0.0478  
Training loss for step 12: 0.2459  
Training loss for step 13: 0.1078  
Training loss for step 14: 0.1651  
Training loss for step 15: 0.0796  
Training loss for step 16: 0.0617  
Training loss for step 17: 0.0187

Epoch 1: Train loss: 0.1137 Validation Loss: 0.0847, Train Accuracy: 0.9722,  
Validation Accuracy 0.9750, Train F1 Score: 0.9583, Validation F1 Score: 0.9600  
Start of epoch 2

Training loss for step 0: 0.1586  
Training loss for step 1: 0.1552  
Training loss for step 2: 0.0875  
Training loss for step 3: 0.0727  
Training loss for step 4: 0.1333  
Training loss for step 5: 0.1190  
Training loss for step 6: 0.0243  
Training loss for step 7: 0.0489  
Training loss for step 8: 0.1300  
Training loss for step 9: 0.0766  
Training loss for step 10: 0.0375  
Training loss for step 11: 0.0951  
Training loss for step 12: 0.0406  
Training loss for step 13: 0.0317  
Training loss for step 14: 0.0338  
Training loss for step 15: 0.0336  
Training loss for step 16: 0.1517  
Training loss for step 17: 0.0118

Epoch 2: Train loss: 0.0801 Validation Loss: 0.0731, Train Accuracy: 0.9722,  
Validation Accuracy 0.9750, Train F1 Score: 0.9581, Validation F1 Score: 0.9600  
Start of epoch 3

Training loss for step 0: 0.0176  
Training loss for step 1: 0.0200  
Training loss for step 2: 0.2108  
Training loss for step 3: 0.0958

Training loss for step 4: 0.0931  
Training loss for step 5: 0.0339  
Training loss for step 6: 0.0630  
Training loss for step 7: 0.1484  
Training loss for step 8: 0.1139  
Training loss for step 9: 0.0108  
Training loss for step 10: 0.0241  
Training loss for step 11: 0.0618  
Training loss for step 12: 0.1229  
Training loss for step 13: 0.0706  
Training loss for step 14: 0.0096  
Training loss for step 15: 0.1235  
Training loss for step 16: 0.0393  
Training loss for step 17: 0.0176

Epoch 3: Train loss: 0.0709 Validation Loss: 0.0706, Train Accuracy: 0.9722,  
Validation Accuracy 0.9812, Train F1 Score: 0.9581, Validation F1 Score: 0.9703  
Start of epoch 4

Training loss for step 0: 0.0235  
Training loss for step 1: 0.0391  
Training loss for step 2: 0.0369  
Training loss for step 3: 0.0217  
Training loss for step 4: 0.0193  
Training loss for step 5: 0.2683  
Training loss for step 6: 0.0583  
Training loss for step 7: 0.0919  
Training loss for step 8: 0.0791  
Training loss for step 9: 0.0810  
Training loss for step 10: 0.0166  
Training loss for step 11: 0.1420  
Training loss for step 12: 0.0626  
Training loss for step 13: 0.0212  
Training loss for step 14: 0.0748  
Training loss for step 15: 0.0520  
Training loss for step 16: 0.0876  
Training loss for step 17: 0.0023

Epoch 4: Train loss: 0.0655 Validation Loss: 0.0708, Train Accuracy: 0.9757,  
Validation Accuracy 0.9812, Train F1 Score: 0.9634, Validation F1 Score: 0.9703

## 1.9 Evaluate the Model


### 1.9.1 Plots for Evaluation

We plot the progress of loss as training proceeds over number of epochs.



```
[39]: def plot_metrics(train_metric, val_metric, metric_name, title, ylim=5):
    plt.title(title)
    plt.ylim(0,ylim)
    plt.gca().xaxis.set_major_locator(mticker.MultipleLocator(1))
    plt.plot(train_metric,color='blue',label=metric_name)
    plt.plot(val_metric,color='green',label='val_' + metric_name)

plot_metrics(epochs_train_losses, epochs_val_losses, "Loss", "Loss", ylim=1.0)
```



We plot the confusion matrix to visualize the true values against the values predicted by the model.

```
[40]: test_outputs = model(norm_test_X.values)
plot_confusion_matrix(test_Y.values, tf.round(test_outputs), title='Confusion_
↳Matrix for Untrained Model')
```

