C2_W4_Lab_2_multi-GPU-mirrored-strategy

July 28, 2021

1 Multi-GPU Mirrored Strategy

In this ungraded lab, you'll go through how to set up a Multi-GPU Mirrored Strategy. The lab environment only has a CPU but we placed the code here in case you want to try this out for yourself in a multiGPU device.

Notes: - If you are running this on Coursera, you'll see it gives a warning about no presence of GPU devices. - If you are running this in Colab, make sure you have selected your `runtime` to be `GPU`. - In both these cases, you'll see there's only 1 device that is available.
- One device is sufficient for helping you understand these distribution strategies.

1.1 Imports

```python
[1]: import tensorflow as tf
     import numpy as np
     import os
```

1.2 Setup Distribution Strategy

```python
[2]: # Note that it generally has a minimum of 8 cores, but if your GPU has
     # less, you need to set this. In this case one of my GPUs has 4 cores
     os.environ["TF_MIN_GPU_MULTIPROCESSOR_COUNT"] = "4"

     # If the list of devices is not specified in the
     # `tf.distribute.MirroredStrategy` constructor, it will be auto-detected.
     # If you have *different* GPUs in your system, you probably have to set up␣
      ↪cross_device_ops like this
     strategy = tf.distribute.MirroredStrategy(cross_device_ops=tf.distribute.
      ↪HierarchicalCopyAllReduce())
     print ('Number of devices: {}'.format(strategy.num_replicas_in_sync))
```

```
INFO:tensorflow:Using MirroredStrategy with devices
('/job:localhost/replica:0/task:0/device:CPU:0',)
Number of devices: 1
```

1.3 Prepare the Data

```python
[3]: # Get the data
fashion_mnist = tf.keras.datasets.fashion_mnist
(train_images, train_labels), (test_images, test_labels) = fashion_mnist.
 ↪load_data()

# Adding a dimension to the array -> new shape == (28, 28, 1)
# We are doing this because the first layer in our model is a convolutional
# layer and it requires a 4D input (batch_size, height, width, channels).
# batch_size dimension will be added later on.
train_images = train_images[..., None]
test_images = test_images[..., None]

# Normalize the images to [0, 1] range.
train_images = train_images / np.float32(255)
test_images = test_images / np.float32(255)

# Batch the input data
BUFFER_SIZE = len(train_images)
BATCH_SIZE_PER_REPLICA = 64
GLOBAL_BATCH_SIZE = BATCH_SIZE_PER_REPLICA * strategy.num_replicas_in_sync

# Create Datasets from the batches
train_dataset = tf.data.Dataset.from_tensor_slices((train_images,␣
 ↪train_labels)).shuffle(BUFFER_SIZE).batch(GLOBAL_BATCH_SIZE)
test_dataset = tf.data.Dataset.from_tensor_slices((test_images, test_labels)).
 ↪batch(GLOBAL_BATCH_SIZE)

# Create Distributed Datasets from the datasets
train_dist_dataset = strategy.experimental_distribute_dataset(train_dataset)
test_dist_dataset = strategy.experimental_distribute_dataset(test_dataset)
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/train-labels-idx1-ubyte.gz
32768/29515 [==============================] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/train-images-idx3-ubyte.gz
26427392/26421880 [==============================] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/t10k-labels-idx1-ubyte.gz
8192/5148 [===================================================] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/t10k-images-idx3-ubyte.gz
4423680/4422102 [==============================] - 0s 0us/step
```

1.4 Define the Model

```python
[4]: # Create the model architecture
     def create_model():
       model = tf.keras.Sequential([
           tf.keras.layers.Conv2D(32, 3, activation='relu'),
           tf.keras.layers.MaxPooling2D(),
           tf.keras.layers.Conv2D(64, 3, activation='relu'),
           tf.keras.layers.MaxPooling2D(),
           tf.keras.layers.Flatten(),
           tf.keras.layers.Dense(64, activation='relu'),
           tf.keras.layers.Dense(10)
         ])
       return model
```

1.5 Configure custom training

Instead of `model.compile()`, we're going to do custom training, so let's do that within a strategy scope.

```python
[5]: with strategy.scope():
         # We will use sparse categorical crossentropy as always. But, instead of
     ↪having the loss function
         # manage the map reduce across GPUs for us, we'll do it ourselves with a
     ↪simple algorithm.
         # Remember -- the map reduce is how the losses get aggregated
         # Set reduction to `none` so we can do the reduction afterwards and divide
     ↪byglobal batch size.
         loss_object = tf.keras.losses.
     ↪SparseCategoricalCrossentropy(from_logits=True, reduction=tf.keras.losses.
     ↪Reduction.NONE)

         def compute_loss(labels, predictions):
             # Compute Loss uses the loss object to compute the loss
             # Notice that per_example_loss will have an entry per GPU
             # so in this case there'll be 2 -- i.e. the loss for each replica
             per_example_loss = loss_object(labels, predictions)
             # You can print it to see it -- you'll get output like this:
             # Tensor("sparse_categorical_crossentropy/weighted_loss/Mul:0",
     ↪shape=(48,), dtype=float32, device=/job:localhost/replica:0/task:0/device:
     ↪GPU:0)
             # Tensor("replica_1/sparse_categorical_crossentropy/weighted_loss/Mul:
     ↪0", shape=(48,), dtype=float32, device=/job:localhost/replica:0/task:0/
     ↪device:GPU:1)
             # Note in particular that replica_0 isn't named in the weighted_loss --
     ↪the first is unnamed, the second is replica_1 etc
```

```
        print(per_example_loss)
        return tf.nn.compute_average_loss(per_example_loss,␣
 ↪global_batch_size=GLOBAL_BATCH_SIZE)


    # We'll just reduce by getting the average of the losses
    test_loss = tf.keras.metrics.Mean(name='test_loss')

    # Accuracy on train and test will be SparseCategoricalAccuracy
    train_accuracy = tf.keras.metrics.
 ↪SparseCategoricalAccuracy(name='train_accuracy')
    test_accuracy = tf.keras.metrics.
 ↪SparseCategoricalAccuracy(name='test_accuracy')

    # Optimizer will be Adam
    optimizer = tf.keras.optimizers.Adam()

    # Create the model within the scope
    model = create_model()
```

1.6   Train and Test Steps Functions

Let's define a few utilities to facilitate the training.

```
[6]: # `run` replicates the provided computation and runs it
     # with the distributed input.
     @tf.function
     def distributed_train_step(dataset_inputs):
       per_replica_losses = strategy.run(train_step, args=(dataset_inputs,))
       #tf.print(per_replica_losses.values)
       return strategy.reduce(tf.distribute.ReduceOp.SUM, per_replica_losses,␣
     ↪axis=None)

     def train_step(inputs):
       images, labels = inputs
       with tf.GradientTape() as tape:
         predictions = model(images, training=True)
         loss = compute_loss(labels, predictions)

       gradients = tape.gradient(loss, model.trainable_variables)
       optimizer.apply_gradients(zip(gradients, model.trainable_variables))

       train_accuracy.update_state(labels, predictions)
       return loss

     #####################
     # Test Steps Functions
```

```
#######################
@tf.function
def distributed_test_step(dataset_inputs):
  return strategy.run(test_step, args=(dataset_inputs,))


def test_step(inputs):
  images, labels = inputs

  predictions = model(images, training=False)
  t_loss = loss_object(labels, predictions)

  test_loss.update_state(t_loss)
  test_accuracy.update_state(labels, predictions)
```

1.7 Training Loop

We can now start training the model.

```
[7]: EPOCHS = 10
     for epoch in range(EPOCHS):
       # Do Training
       total_loss = 0.0
       num_batches = 0
       for batch in train_dist_dataset:
         total_loss += distributed_train_step(batch)
         num_batches += 1
       train_loss = total_loss / num_batches

       # Do Testing
       for batch in test_dist_dataset:
         distributed_test_step(batch)

       template = ("Epoch {}, Loss: {}, Accuracy: {}, Test Loss: {}, " "Test␣
     ↪Accuracy: {}")

       print (template.format(epoch+1, train_loss, train_accuracy.result()*100,␣
     ↪test_loss.result(), test_accuracy.result()*100))

       test_loss.reset_states()
       train_accuracy.reset_states()
       test_accuracy.reset_states()
```

WARNING:tensorflow:From /opt/conda/lib/python3.7/site-
packages/tensorflow/python/data/ops/multi_device_iterator_ops.py:601:
get_next_as_optional (from tensorflow.python.data.ops.iterator_ops) is
deprecated and will be removed in a future version.

```
Instructions for updating:
Use `tf.data.Iterator.get_next_as_optional()` instead.
Tensor("sparse_categorical_crossentropy/weighted_loss/Mul:0", shape=(64,),
dtype=float32, device=/job:localhost/replica:0/task:0/device:CPU:0)
Tensor("sparse_categorical_crossentropy/weighted_loss/Mul:0", shape=(64,),
dtype=float32, device=/job:localhost/replica:0/task:0/device:CPU:0)
Tensor("sparse_categorical_crossentropy/weighted_loss/Mul:0", shape=(32,),
dtype=float32, device=/job:localhost/replica:0/task:0/device:CPU:0)
Epoch 1, Loss: 0.513083815574646, Accuracy: 81.30166625976562, Test Loss:
0.38477808237075806, Test Accuracy: 86.22999572753906
Epoch 2, Loss: 0.33968618512153625, Accuracy: 87.62166595458984, Test Loss:
0.3397756814956665, Test Accuracy: 87.6500015258789
Epoch 3, Loss: 0.2946312129497528, Accuracy: 89.23332977294922, Test Loss:
0.3231080174446106, Test Accuracy: 88.25
Epoch 4, Loss: 0.26395100355148315, Accuracy: 90.30833435058594, Test Loss:
0.30268099904060364, Test Accuracy: 88.97000122070312
Epoch 5, Loss: 0.24057503044605255, Accuracy: 91.16999816894531, Test Loss:
0.2736639380455017, Test Accuracy: 89.9000015258789
Epoch 6, Loss: 0.22174939513206482, Accuracy: 91.84166717529297, Test Loss:
0.2691769599914551, Test Accuracy: 90.20999908447266
Epoch 7, Loss: 0.20393218100070953, Accuracy: 92.50666809082031, Test Loss:
0.25806760787963867, Test Accuracy: 90.58000183105469
Epoch 8, Loss: 0.18638382852077484, Accuracy: 93.1883316040039, Test Loss:
0.260176420211792, Test Accuracy: 90.58999633789062
Epoch 9, Loss: 0.17518101632595062, Accuracy: 93.57167053222656, Test Loss:
0.26173508167266846, Test Accuracy: 90.95999908447266
Epoch 10, Loss: 0.16134648025035858, Accuracy: 93.9949951171875, Test Loss:
0.2608034908771515, Test Accuracy: 90.80999755859375
```

[ ]: